

Iterative Functional Reactive Programming with the Nu Game Engine

An Informal Experience Report

NOTE: This is an entirely informal 'paper', originally targeted for a simple post on a forum of <http://lambda-the-ultimate.org>, but due to my inability to achieve a readable formatting with the tools of said forum, is made available via this PDF.

So, I have this purely functional game engine in F# - <https://github.com/bryanedds/FPWorks>

For readers who wish to build some familiarity with the system before reading this admittedly terse description, please look here for Nu's high-level documentation - <https://github.com/bryanedds/FPWorks/blob/master/Nu/Documentation/Nu%20Game%20Engine.pdf?raw=true>

The Nu Game Engine certainly qualifies as 'functional reactive' in that -

- 1) it is reactive in that it uses user-defined events to derive successive game states.
- 2) it is functional in that it uses pure functions everywhere, even in the event system.

However, I cannot describe it as the typical first-order or higher-order FRP system since it uses neither continuous nor discrete functions explicitly parameterized with time. Instead it uses a classically-iterative approach to advancing game states that is akin to the imperative tick-based style, but implemented with pure functions.

Thus, I gave it the name of 'Iterative FRP'. It's a purely-functional half-way step between classic imperative game programming and first / higher-order FRP systems. There are plusses and minuses to using this 'Iterative FRP' style -

In the plus column -

- 1) it's perfectly straight-forward to implement most any arbitrary game behavior without worrying about things like space or time leaks, or having to invoke higher-ordered forms of expression.
- 2) end-users don't need to understand new forms of expression beyond primitive functional expressions to encode their game logic. No lifting is required (though it is available).
- 3) it has a level of pluggability, data-drivenness, and serializability that are extremely difficult to pull off in FO/HOFRP systems.
- 4) none of its design is the subject of open research questions, and is therefore simple enough for non-PhD's to understand as well as complete in its implementation.

In the minus column -

- 1) debuggability suffers from the declarative nature of the API. For example, - **TODO**
- 2) verbosity - **TODO**

Despite these minuses, I currently conclude IFRP is appropriate for such a general-purpose game engine because -

- 1) from my experience, games, and game technology in general, refuse to conform to any single system of thought, often demanding certain features be implemented with escape-hatch approaches like encapsulated mutation in order to be efficient, or with lower-level forms of functional expression for flexibility.
- 2) just because the default means of expression is not as abstract as that of FO/HOFRP systems, it does not mean there are no available ways to 'climb into' more abstract forms of expression.

To expand on point 2, here are some of the higher forms of expression provided by the API -

Here we have an expression form called an 'observation' that allows combinations and transformations of events like so -

```
let observation =
  observe (ClickEventAddress ->>- hudHalt.EntityAddress) address |>
  filter isSelected |>
  sum TickEventAddress |>
  until (DeselectEventAddress ->>- gameplay.ScreenAddress)
```

This affords us the ability to treat events like first-class collections.

Here we have an expression form called a 'chain'. A chain is a monad that allows a procedural-style expression to span 0 or more events while also taking the world as an implicit state -

```
let chain = chain {
  do! update ^ character.SetActivityState ^ Action newActionDescriptor
  do! during (fun world -> ActivityState.isActing <| character.GetActivityState world) ^ chain {
    do! update ^ fun world ->
      let actionDescriptor =
        match character.GetActivityState world with
        | Action actionDescriptor -> actionDescriptor
        | _ -> failwithmf ()
      let world = updateCharacterByAction actionDescriptor character world
      runCharacterReaction actionDescriptor character gameplay world
    do! pass }}
}
```

We also have the ability to compose chains overs observations like so -

```
let observation = observe TickEventAddress character |> until (DeselectEventAddress ->>- gameplay.ScreenAddress)
snd <| runAssumingCascade chain observation world
```

Finally, there is a means to declaratively forward the changing value of a simulant's field to that of another -

```
let world =
  world |>
  (bob, bob.GetVisible) *--> (jim, jim.SetVisible) |>
  (jim, jim.GetVisible) /--> (bob, not >> bob.SetVisible)
```

Here, `*-->` denotes the forwarding of the value of Bob's Visible field to Jim's Visible field. To throw a monkey-wrench into the declaration, `/-->` is used to in turn forward the value of Jim's Visible field back to Bob's Visible field, albeit **not**'d and with cycle-breaking so that the circularity of this expression is broken appropriately. This is a contrived example, but illustrative of the API's expressiveness.

So, depending on the nature of the game behavior you want to implement, the API provides enough surface area to do things at multiple points of a spectrum of efficiency and flexibility. But there is a down-side to an API with a large surface area... With multiple expression forms and levels of abstractions at which to operate, the learning curve can be regrettably steepened.

Even with functional programming, and even assuming my approach was optimal for the given point in my targeted design space, game engine design remains fundamentally difficult and peppered with compromise.